


Problem 1. Let D be a point outside $\odot ABC$ (circumcircle of triangle ABC) such that points B and D fall on different sides of the line passing through segment \overline{AC} . Then, $\odot ACD$ covers territory of arc \widehat{AC} inside $\odot ABC$ (the shadow region in the figure below).



Hint: Pick any point E on the arc \widehat{AC} ($E \neq A$ and $E \neq C$). Prove that the angle $\angle AEC > \angle ADC$.

Proof. Pick any point E on the arc \widehat{AC} .

Consider the circumcircle $\odot ABC$. Since D is a point outside $\odot ABC$, $\angle AEC > \angle ADC$.

Now consider the circumcircle $\odot ACD$. Now, D is a point on $\odot ACD$. Since $\angle AEC > \angle ADC$, E is inside the circle $\odot ACD$.

Thus, the entire arc \widehat{AC} is covered by $\odot ACD$. □

Problem 2. Given a triangle ABC and a point p , determine in $O(1)$ time if $\odot ABC$ covers p .

Solution. Let $A = (x_1, y_1)$, $B = (x_2, y_2)$, $C = (x_3, y_3)$. Let $q = (x, y)$ be a point on or in the interior of the circumcircle of $\triangle ABC$.

Consider the set of equations:

$$\begin{cases} a(x^2 + y^2) + bx + cy + d \leq 0 \\ a(x_1^2 + y_1^2) + bx_1 + cy_1 + d = 0 \\ a(x_2^2 + y_2^2) + bx_2 + cy_2 + d = 0 \\ a(x_3^2 + y_3^2) + bx_3 + cy_3 + d = 0 \end{cases}$$

To obtain a valid non-trivial solution (a, b, c, d) ,

$$\begin{vmatrix} x^2 + y^2 & x & y & 1 \\ x_1^2 + y_1^2 & x_1 & y_1 & 1 \\ x_2^2 + y_2^2 & x_2 & y_2 & 1 \\ x_3^2 + y_3^2 & x_3 & y_3 & 1 \end{vmatrix} \leq 0.$$

This determinant can be computed in $O(1)$ time.

Problem 3. Let P be a set of n points in \mathbb{R}^2 . A Euclidean spanning tree of P is a tree where every vertex is a point in P and every edge is a line segment connecting two points of P . The tree's weight equals the total (Euclidean) length of all the segments. The *Euclidean minimum spanning tree* (EMST) is a Euclidean spanning tree of the smallest weight. Give an algorithm to find an EMST of P in $O(n \log n)$ expected time.


Remark: The expected $O(n \log n)$ time is due to the randomized Delaunay triangulation algorithm. Deterministic $O(n \log n)$ time algorithm exists.

Solution. First, we compute the point set Delaunay triangulation in expected $O(n \log n)$ time. Then compute the MST of the triangulation by Kruskal's algorithm and return the result.

Note that there are at most $3n-6$ edges in a point set triangulation. Therefore, Kruskal's algorithm can be done in $O(n \log n)$ time.

We now prove that the MST of P is a subgraph of the Delaunay triangulation.

Let T be the MST for P , let $w(T)$ denote the total weight of T . Let a and b be any two sites such that ab is an edge of T . Suppose to the contrary that ab is not an edge in the Delaunay triangulation. This implies that there is no empty circle passing through a and b , and in particular, the circle whose diameter is the segment \overline{ab} contains a site, call it c .



The removal of \overline{ab} from the MST splits the tree into two subtrees. Assume without loss of generality that c lies in the same subtree as a . Now remove the edge \overline{ab} from the MST and add the edge \overline{bc} in its place. This will incur in a spanning tree T' whose weight is

$$w(T') = w(T) + |bc| - |ab| < w(T).$$

Since ab is the diameter of the circle, $|bc| < |ab|$. This contradicts the hypothesis that T is the MST, completing the proof.

Problem 4. Let P be a set of n points in \mathbb{R}^2 . A tour is a sequence of n segments $\overline{p_1 p_2}, \overline{p_2 p_3}, \dots, \overline{p_{n-1} p_n}, \overline{p_n p_1}$, where each $p_i (i \in [1, n])$ is a distinct point in P . The *length* of the tour is the total length of all the n segments. Let ℓ^* be the shortest length of all possible tours. Design an algorithm to find a tour with length at most $2\ell^*$ in $O(n \log n)$ expected time.

Solution. In $O(n \log n)$ expected time, one can obtain the MST T of P . Then, a closed walk π where all edges on T appear twice can be computed by pre-order traversal in $O(n)$ time. Then, a Hamiltonian cycle σ can be found by scanning $\pi = (\pi_1, \pi_2, \dots, \pi_{2n-1})$. Namely, add π_i to σ if π_i has not been visited before, also add π_1 again to the end of the sequence.

Lemma 1: $\ell^* \geq \text{MST}(T)$. Consider removing an edge e from the cycle attaining the cost ℓ^* , we obtained a spanning tree of P . By no means that $\ell^* - \text{cost}(e) < \text{MST}(T)$ and thus $\ell^* \geq \text{MST}(T)$.

Lemma 2: $\left(\sum_{i=1}^{n-1} \text{dist}(\sigma_i, \sigma_{i+1}) \right) + \text{dist}(\sigma_n, \sigma_1) \leq 2 \cdot \text{MST}(T)$.

Let $\sigma_1 = \pi_{j_1}, \sigma_2 = \pi_{j_2}, \dots, \sigma_n = \pi_{j_n}$. It is apparent that $1 = j_1 < j_2 < \dots < j_n = 2n - 1$.

Therefore by the triangle inequality,

$$\text{dist}(\sigma_i, \sigma_{i+1}) \leq \sum_{k=j_i}^{j_{i+1}-1} \text{dist}(\pi_k, \pi_{k+1}),$$

meaning that the length of cycle $\sigma \leq$ length of cycle $\pi = 2 \cdot \text{MST}(T) \leq 2\ell^*$.

The sequence σ is what we want to find.

Problem 5. Let P be a set of n points in \mathbb{R}^2 . A k -clustering of P is a partition P_1, P_2, \dots, P_k of P such that

- each $P_i (i \in [1, k])$ is a non-empty subset of P ,
- $P_i \cap P_j = \emptyset$ for any different $i, j \in [1, k]$,
- $P_1 \cup P_2 \cup P_3 \cup \dots \cup P_k = P$.

We will refer to each $P_i (i \in [1, k])$ as a cluster. For any different $i, j \in [1, k]$, define the distance between clusters P_i and P_j as

$$\text{dist}(P_i, P_j) = \min_{p \in P_i, q \in P_j} \text{dist}(p, q)$$

where $\text{dist}(p, q)$ is the Euclidean distance between p and q . The quality of the k -clustering is defined to be the smallest distance between all $\binom{k}{2}$ cluster pairs.

- Prove: If the quality of P is determined by two points $p, q \in P$, then \overline{pq} is an edge in the Delaunay triangulation of P .
- Given a real value $r > 0$ and an integer $k \geq 2$, give an algorithm to determine whether there exists a k -clustering of P whose quality is at least r . Your algorithm needs to finish in $O(n \log n)$ expected time.

Hint: The first question is easy. For the second question, what happens if we remove all the edges of the Delaunay graph of P that have lengths less than r ?

Solution.

- Suppose the otherwise that \overline{pq} is not an edge in the Delaunay triangulation of P . There must exist a point r lying inside the circle having \overline{pq} as the diameter. Let o be the center of this circle. Then, without loss of generality, suppose that r does not belong to the cluster with p , $|op| + |or| < 2R = |op| + |oq| = |pq|$. This is a contradiction as the cross-pair distance of the given clustering is at least $|pq|$.
- We first spend $O(n \log n)$ time to compute the Delaunay triangulation of P . Then, we can sort the list of edges L in ascending order in the triangulation in $O(n \log n)$ time. Start by having every point as disjoint n clusters. Now, consider each edge $(u, v) \in L$, we merge the two clusters with u and v only if they do not belong to the same cluster. We terminate the process,
 - if we have to merge more than $k - 1$ times, and report “no”;
 - otherwise proceed until we have merged $k - 1$ times, and report “yes”.

It is easy to see the algorithm’s correctness as

1. If we have reported “yes”, no quality less than r is possible as each point pair connected by such edges are in the same cluster now.
2. If we have reported “no”, there must exist an edge (u, v) with $\text{dist}(u, v) < r$ which fails to be merged. It then means that u and v is a cross-cluster pair, the quality must count this pair in, and is upper-bounded by the distance of this pair.